

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
df = pd.read_csv(r"C:\Users\chava\Desktop\
heart_disease_prediction.csv")
df
```

	Age	Sex	ChestPainType	RestingBP	Cholesterol	RestingECG	MaxHR	\
0	40	M	ATA	140	289	Normal	172	
1	49	F	NAP	160	180	Normal	156	
2	37	M	ATA	130	283	ST	98	
3	48	F	ASY	138	214	Normal	108	
4	54	M	NAP	150	195	Normal	122	
...	
913	45	M	TA	110	264	Normal	132	
914	68	M	ASY	144	193	Normal	141	
915	57	M	ASY	130	131	Normal	115	
916	57	F	ATA	130	236	LVH	174	
917	38	M	NAP	138	175	Normal	173	

	ExerciseAngina	Oldpeak	ST_Slope	HeartDisease
0	No	0.0	Up	Absence
1	No	1.0	Flat	Presence
2	No	0.0	Up	Absence
3	Yes	1.5	Flat	Presence
4	No	0.0	Up	Absence
...
913	No	1.2	Flat	Presence
914	No	3.4	Flat	Presence
915	Yes	1.2	Flat	Presence
916	No	0.0	Flat	Presence
917	No	0.0	Up	Absence

[918 rows x 11 columns]

EDA process part in Dataset

```
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 918 entries, 0 to 917
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Age                    918 non-null    int64
```

```

1  Sex          918 non-null  object
2  ChestPainType 918 non-null  object
3  RestingBP     918 non-null  object
4  Cholesterol   918 non-null  object
5  RestingECG    918 non-null  object
6  MaxHR         918 non-null  int64
7  ExerciseAngina 918 non-null  object
8  Oldpeak       918 non-null  float64
9  ST_Slope      918 non-null  object
10 HeartDisease  918 non-null  object
dtypes: float64(1), int64(2), object(8)
memory usage: 79.0+ KB

```

need to count target column values, for our data set good or not

```
df['HeartDisease'].value_counts()
```

```

Presence    508
Absence     410
Name: HeartDisease, dtype: int64

```

i found ? in RestingBP and Cholesterol column, so i replace the ? with NAN because pandas and sklearn only handle NaN values

```

df["RestingBP"].replace("?", np.nan, inplace=True)
df["Cholesterol"].replace("?", np.nan, inplace=True)

```

change the datatype because in the latest steps we need to get mean value of RestingBP and Cholesterol column

```

df["RestingBP"] = df["RestingBP"].astype("float64")
df["Cholesterol"] = df["Cholesterol"].astype("float64")

df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 918 entries, 0 to 917
Data columns (total 11 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Age             918 non-null   int64
1   Sex             918 non-null   object
2   ChestPainType   918 non-null   object
3   RestingBP       910 non-null   float64
4   Cholesterol      911 non-null   float64
5   RestingECG      918 non-null   object
6   MaxHR           918 non-null   int64
7   ExerciseAngina  918 non-null   object
8   Oldpeak         918 non-null   float64
9   ST_Slope        918 non-null   object


```

```
10 HeartDisease    918 non-null    object
dtypes: float64(3), int64(2), object(6)
memory usage: 79.0+ KB
```

```
df.describe()
```


	Age	RestingBP	Cholesterol	MaxHR	Oldpeak
count	918.000000	910.000000	911.000000	918.000000	918.000000
mean	53.510893	132.400000	198.656422	136.809368	0.887364
std	9.432617	18.562723	109.753487	25.460334	1.066570
min	28.000000	0.000000	0.000000	60.000000	-2.600000
25%	47.000000	120.000000	172.500000	120.000000	0.000000
50%	54.000000	130.000000	223.000000	138.000000	0.600000
75%	60.000000	140.000000	267.000000	156.000000	1.500000
max	77.000000	200.000000	603.000000	202.000000	6.200000

```
plt.grid()
sns.boxplot(df["RestingBP"])
<AxesSubplot:xlabel='RestingBP'>
```



from describe i found that RestingBP has min 0 value, but RestingBP must be greater than zero,so i need to 1st replace 0 value with NAN and then again NAN repalce with mean of RestingBP

```
df["RestingBP"].replace(0, np.nan, inplace=True)
df["RestingBP"].replace(np.nan,df["RestingBP"].mean(),inplace=True)
plt.grid()
sns.boxplot(df["Cholesterol"])
<AxesSubplot:xlabel='Cholesterol'>
```



```
df[(df['Cholesterol']==0)]
```

	Age	Sex	ChestPainType	RestingBP	Cholesterol	RestingECG	
MaxHR \							
293	65	M	ASY	115.0	0.0	Normal	93
294	32	M	TA	95.0	0.0	Normal	127
295	61	M	ASY	105.0	0.0	Normal	110
296	50	M	ASY	145.0	0.0	Normal	139

297	57	M	ASY	110.0	0.0	ST	131
..
514	43	M	ASY	122.0	0.0	Normal	120
515	63	M	NAP	130.0	0.0	ST	160
518	48	M	NAP	102.0	0.0	ST	110
535	56	M	ASY	130.0	0.0	LVH	122
536	62	M	NAP	133.0	0.0	ST	119

	ExerciseAngina	Oldpeak	ST_Slope	HeartDisease
293	Yes	0.0	Flat	Presence
294	No	0.7	Up	Presence
295	Yes	1.5	Up	Presence
296	Yes	0.7	Flat	Presence
297	Yes	1.4	Up	Presence
..
514	No	0.5	Up	Presence
515	No	3.0	Flat	Absence
518	Yes	1.0	Down	Presence
535	Yes	1.0	Flat	Presence
536	Yes	1.2	Flat	Presence

[172 rows x 11 columns]

in Cholesterol i found that 172 rows 0 value,Cholesterol must be greater than zero

so 1st i need to replace this 0 value with NAN and then replace NAN value with Cholesterol mean value

```
df["Cholesterol"].replace(0, np.nan, inplace=True)
df["Cholesterol"].replace(np.nan,df["Cholesterol"].mean(),inplace=True)
df.describe()
```

	Age	RestingBP	Cholesterol	MaxHR	Oldpeak
count	918.000000	918.000000	918.000000	918.000000	918.000000
mean	53.510893	132.545655	244.893099	136.809368	0.887364
std	9.432617	17.956368	53.186105	25.460334	1.066570
min	28.000000	80.000000	85.000000	60.000000	-2.600000
25%	47.000000	120.000000	215.000000	120.000000	0.000000
50%	54.000000	130.000000	244.893099	138.000000	0.600000
75%	60.000000	140.000000	267.000000	156.000000	1.500000
max	77.000000	200.000000	603.000000	202.000000	6.200000

in oldpeak column positive and negative value between -2.60 to 6.20, so we need to scaling this column

```
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
df[['Oldpeak']] = scaler.fit_transform(df[['Oldpeak']])
df.describe()
```

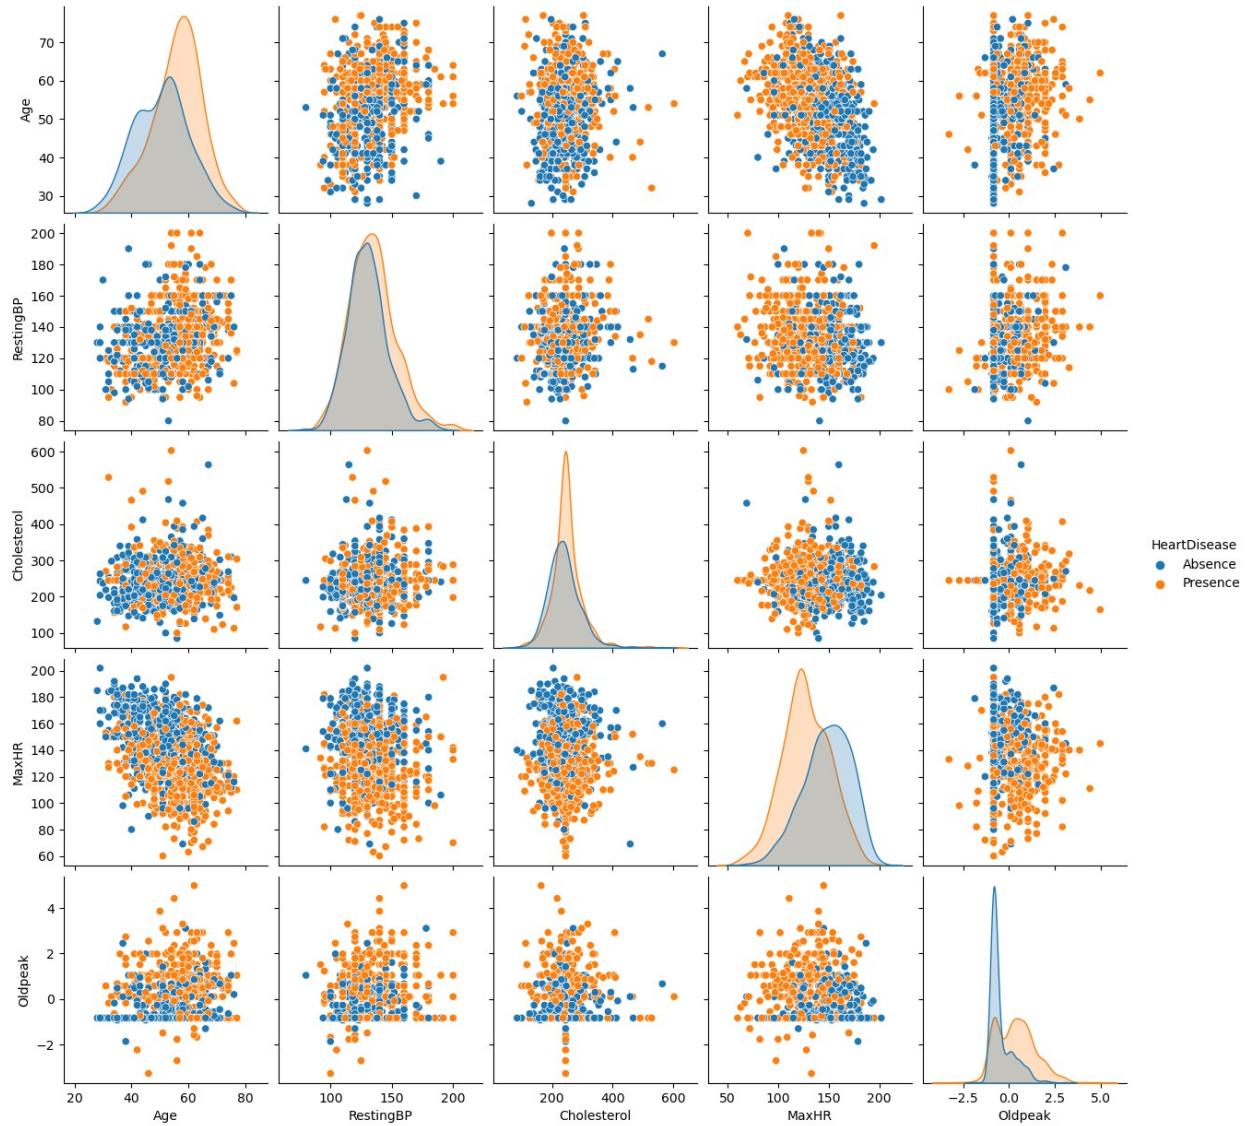
	Age	RestingBP	Cholesterol	MaxHR	Oldpeak
count	918.000000	918.000000	918.000000	918.000000	9.180000e+02
mean	53.510893	132.545655	244.893099	136.809368	-2.024524e-16
std	9.432617	17.956368	53.186105	25.460334	1.000545e+00
min	28.000000	80.000000	85.000000	60.000000	-3.271482e+00
25%	47.000000	120.000000	215.000000	120.000000	-8.324324e-01
50%	54.000000	130.000000	244.893099	138.000000	-2.695748e-01
75%	60.000000	140.000000	267.000000	156.000000	5.747115e-01
max	77.000000	200.000000	603.000000	202.000000	4.983762e+00

now our dataframe EDA part finished, so we can start data visualization

data visualization in dataset

data visualization for numerical dataset

```
sns.pairplot(df[['Age', 'RestingBP', 'Cholesterol', 'MaxHR',
'Oldpeak', 'HeartDisease']], hue='HeartDisease')
plt.show()
```





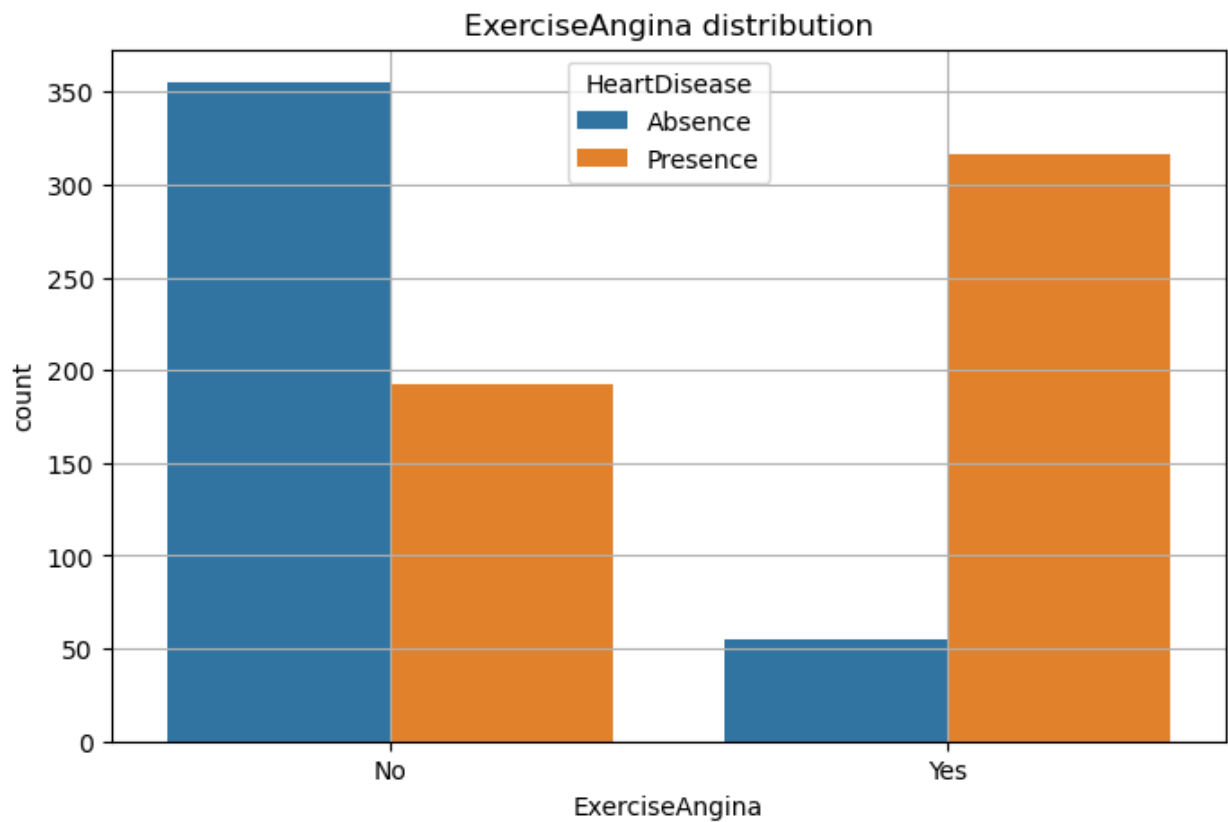
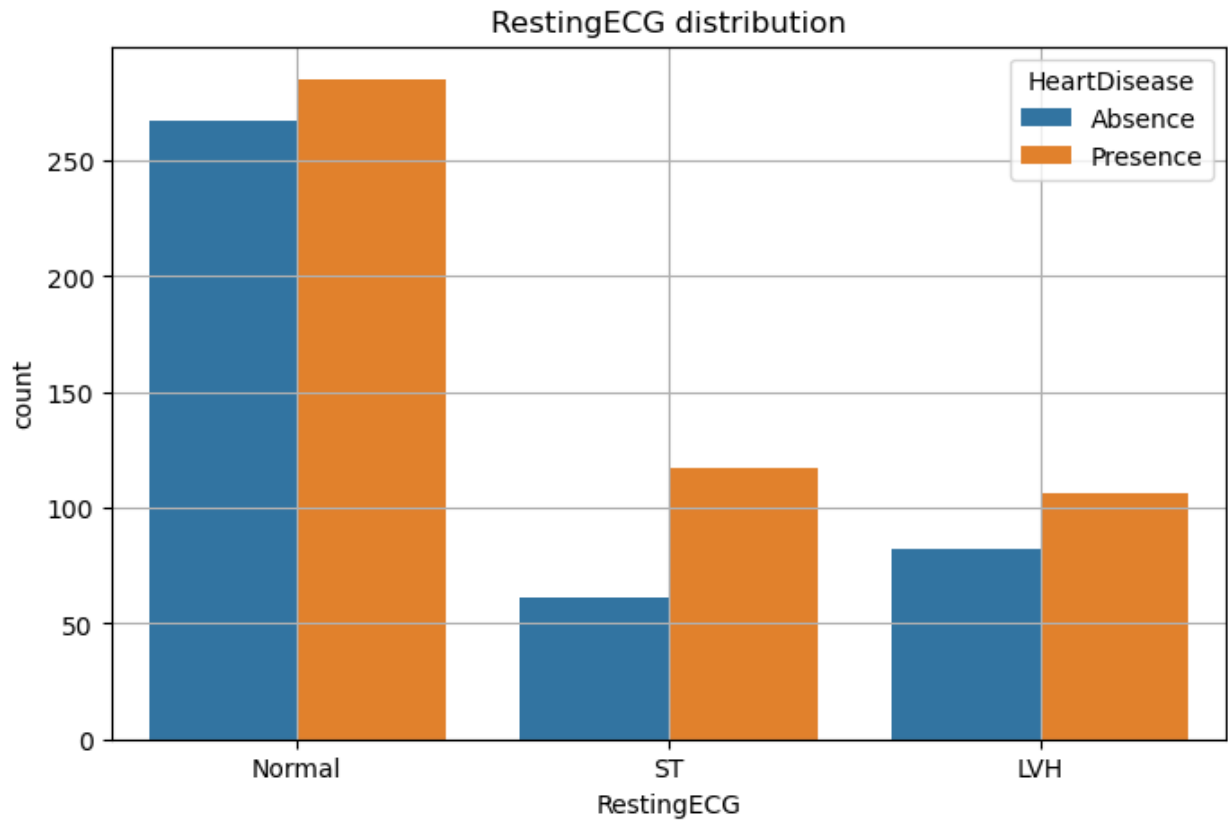
data visualization for categorical dataset

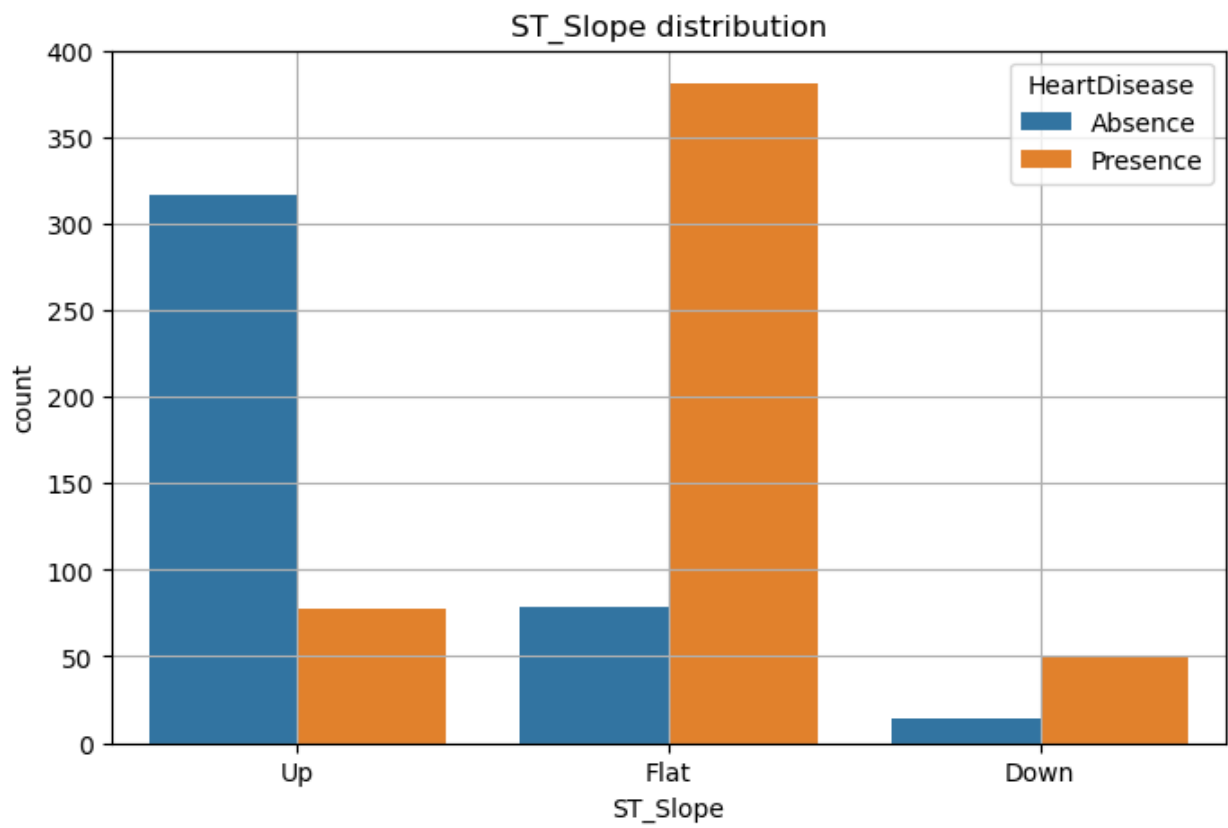
```

categorical_columns = ['Sex', 'ChestPainType', 'RestingECG',
'ExerciseAngina', 'ST_Slope']
for column in categorical_columns:
    plt.figure(figsize=(8, 5))
    sns.countplot(x=column, hue='HeartDisease', data=df)
    plt.title(f'{column} distribution')
    plt.grid()
    plt.show()

```







(1) from above 1st male and female bar chart i can conclusion that male Heart disease presence 63% and female 25% meance heart disease presence percentage in male greater than female

(2) from above 2nd bar chart ChestPainType i can conclusion that Heart disease presence in ASY (asymptomatic) 79%,ATA (atypical angina) 13%,NAP(non-anginal pain) 35%,TA(truncus arteriosus) 43% meance ASY most important role in heart disease presence

(3)above 4th bar graph ExerciseAngina i can conclusion that heart disease presence due to doing exercise 85% and not doing exercise 35% meance due to Exercise heart disease increase

(4) above 4th bar graph ST Slope i can conclusion that Heart disease presence in up 19% ,flat 82% and down 77% meance the ST segment shift relative to exercise-induced increments in heart rate and flat slope increase heart disease presence percentage

Machine Learning with Logistic Regression

import important libraries for logistic regression

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix
from sklearn.preprocessing import StandardScaler, LabelEncoder
```

Preprocessing: Convert categorical variables to numerical

```
le = LabelEncoder()
df['Sex'] = le.fit_transform(df['Sex'])
df['ChestPainType'] = le.fit_transform(df['ChestPainType'])
df['RestingECG'] = le.fit_transform(df['RestingECG'])
df['ExerciseAngina'] = le.fit_transform(df['ExerciseAngina'])
df['ST_Slope'] = le.fit_transform(df['ST_Slope'])
```

Separate features (X) and target variable (y)

```
X = df.drop('HeartDisease', axis=1)
y = df['HeartDisease']
```

Split the data into training and testing sets

```
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
```

Initialize and train the logistic regression model

```
model = LogisticRegression()  
model.fit(X_train, y_train)
```

```
LogisticRegression()
```

Make predictions on the test set

```
y_pred = model.predict(X_test)
```

Evaluate the model

```
accuracy = accuracy_score(y_test, y_pred)  
conf_matrix = confusion_matrix(y_test, y_pred)  
classification_rep = classification_report(y_test, y_pred)
```

```
print(f'Accuracy: {accuracy}')
```

```
print(f'Confusion Matrix:\n{conf_matrix}')
```

```
print(f'Classification Report:\n{classification_rep}')
```

Accuracy: 0.8260869565217391

Confusion Matrix:

```
[[68  9]  
 [23 84]]
```

Classification Report:

	precision	recall	f1-score	support
Absence	0.75	0.88	0.81	77
Presence	0.90	0.79	0.84	107
accuracy			0.83	184
macro avg	0.83	0.83	0.82	184
weighted avg	0.84	0.83	0.83	184

Machine Learning with DecisionTreeClassifier

import important libraries for DecisionTreeClassifier

```
import pandas as pd  
from sklearn.model_selection import train_test_split  
from sklearn.tree import DecisionTreeClassifier  
from sklearn.metrics import accuracy_score, classification_report,  
confusion_matrix  
from sklearn.preprocessing import StandardScaler, LabelEncoder
```

Preprocessing: Convert categorical variables to numerical

```
le = LabelEncoder()  
df['Sex'] = le.fit_transform(df['Sex'])  
df['ChestPainType'] = le.fit_transform(df['ChestPainType'])  
df['RestingECG'] = le.fit_transform(df['RestingECG'])  
df['ExerciseAngina'] = le.fit_transform(df['ExerciseAngina'])  
df['ST_Slope'] = le.fit_transform(df['ST_Slope'])
```

Separate features (X) and target variable (y)

```
X = df.drop('HeartDisease', axis=1)  
y = df['HeartDisease']
```

Split the data into training and testing sets

```
X_train, X_test, y_train, y_test = train_test_split(X, y,  
test_size=0.2, random_state=42)
```

Initialize and train the decision tree model

```
model = DecisionTreeClassifier(random_state=42)  
model.fit(X_train, y_train)  
model = DecisionTreeClassifier(max_depth=5, min_samples_leaf=5,  
random_state=42)  
model.fit(X_train, y_train)
```

```
DecisionTreeClassifier(max_depth=5, min_samples_leaf=5,  
random_state=42)
```

Make predictions on the test set

```
y_pred = model.predict(X_test)
```

Evaluate the model

```
accuracy = accuracy_score(y_test, y_pred)  
conf_matrix = confusion_matrix(y_test, y_pred)  
classification_rep = classification_report(y_test, y_pred)
```

```
print(f'Accuracy: {accuracy}')
```

```
print(f'Confusion Matrix:\n{conf_matrix}')
```

```
print(f'Classification Report:\n{classification_rep}')
```

Accuracy: 0.8478260869565217

Confusion Matrix:

```
[[69  8]
```

```
 [20 87]]
```

Classification Report:

	precision	recall	f1-score	support
Absence	0.78	0.90	0.83	77

Presence	0.92	0.81	0.86	107
accuracy			0.85	184
macro avg	0.85	0.85	0.85	184
weighted avg	0.86	0.85	0.85	184

Machine Learning with SVC

import important libraries for SVC

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix
from sklearn.preprocessing import StandardScaler, LabelEncoder
```

Preprocessing: Convert categorical variables to numerical

```
le = LabelEncoder()
df['Sex'] = le.fit_transform(df['Sex'])
df['ChestPainType'] = le.fit_transform(df['ChestPainType'])
df['RestingECG'] = le.fit_transform(df['RestingECG'])
df['ExerciseAngina'] = le.fit_transform(df['ExerciseAngina'])
df['ST_Slope'] = le.fit_transform(df['ST_Slope'])
```

Separate features (X) and target variable (y)

```
X = df.drop('HeartDisease', axis=1)
y = df['HeartDisease']
```

Split the data into training and testing sets

```
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
```

Standardize the features (SVMs are sensitive to feature scaling)

```
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

Initialize and train the SVM model

```
model = SVC(C=1.0, gamma='scale', random_state=42)
model.fit(X_train, y_train)

SVC(random_state=42)
```

Make predictions on the test set

```
y_pred = model.predict(X_test)
```

Evaluate the model

```
accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
classification_rep = classification_report(y_test, y_pred)
```

```
print(f'Accuracy: {accuracy}')
print(f'Confusion Matrix:\n{conf_matrix}')
print(f'Classification Report:\n{classification_rep}')
```

Accuracy: 0.8369565217391305

Confusion Matrix:

```
[[64 13]
 [17 90]]
```

Classification Report:

	precision	recall	f1-score	support
Absence	0.79	0.83	0.81	77
Presence	0.87	0.84	0.86	107
accuracy			0.84	184
macro avg	0.83	0.84	0.83	184
weighted avg	0.84	0.84	0.84	184